

SocialNLP EmotionX 2019 Challenge Overview: Predicting Emotions in Spoken Dialogues and Chats

Boaz Shmueli^{1, 2*}, Lun-Wei Ku³

¹Social Networks and Human-Centered Computing, Taiwan International Graduate Program (TIGP)

Institute of Information Science, Academia Sinica

²National Tsing Hua University

³Institute of Information Science, Academia Sinica
{shmueli, lwku}@iis.sinica.edu.tw

Abstract

We present an overview of the EmotionX 2019 Challenge, held at the 7th International Workshop on Natural Language Processing for Social Media (SocialNLP), in conjunction with IJCAI 2019. The challenge entailed predicting emotions in spoken and chat-based dialogues using augmented EmotionLines datasets. EmotionLines contains two distinct datasets: the first includes excerpts from a US-based TV sitcom episode scripts (Friends) and the second contains real online chats (EmotionPush). A total of thirty-six teams registered to participate in the challenge. Eleven of the teams successfully submitted their predictions for performance evaluation. The top-scoring team achieved a micro-F1 score of 81.5% for the spoken-based dialogues (Friends) and 79.5% for the chat-based dialogues (EmotionPush).

1 Introduction

Emotions are a central component of our existence as human beings, and are manifested by physiological and psychological changes that often affect behavior and action. Emotions involve a complicated interplay of mind, body, language, and culture [Bazzanella, 2004].

Detecting and recognizing emotions is a difficult task for machines. Nevertheless, following the successful use of computational linguistics to analyze sentiment in texts, there is growing interest in the more difficult task of the automatic detection and classification of emotions in texts.

The detection of emotions in text is a complicated challenge for multiple reasons: first, emotions are complex entities, and no universally-agreed upon psychological model of emotions exists. Second, isolated texts convey less information compared to a complete human interaction in which emotions can be detected from the other person’s facial expressions, listening to their tone of voice, etc. However, due to important applications in fields such as psychology, marketing, and political science, research in this topic is now expanding rapidly [Seyeditabari *et al.*, 2018].

In particular, dialogue systems such as those available on social media or instant messaging services are rich sources of textual data and have become the focus of much attention. Emotions of utterances within dialogues can be detected more precisely due to the presence of more context. For example, a single utterance (“OK!”) might convey different emotions (happiness, anger, surprise), depending on its context. Taking all this into consideration, in 2018 the EmotionX Challenge asked participants to detect emotions in complete dialogues [Hsu and Ku, 2018]. Participants were challenged to classify utterances using Ekman’s well-known theory of six basic emotions (sadness, happiness, anger, fear, disgust, and surprise) [Ekman, 1999].

For the 2019 challenge, we built and expanded upon the 2018 challenge. We provided an additional 20% of data for training, as well as augmenting the dataset using two-way translation. The metric used was micro-F1 score, and we also report the macro-F1 score.

A total of thirty-six teams registered to participate in the challenge. Eleven of the teams successfully submitted their data for performance evaluation, and seven of them submitted technical papers for the workshop. Approaches used by the teams included deep neural networks and SVM classifiers. In the following sections we expand on the challenge and the data. We then briefly describe the various approaches used by the teams, and conclude with a summary and some notes. Detailed descriptions of the various submissions are available in the teams’ technical reports.

2 Datasets

The two datasets used for the challenge are Friends and EmotionPush, part of the EmotionLines corpus [Chen *et al.*, 2018]. The datasets contain English-language dialogues of varying lengths. For the competition, we provided 1,000 labeled dialogues from each dataset for training, and 240 unlabeled dialogues from each dataset for evaluation. The Friends dialogues are scripts taken from the American TV sitcom (1994-2004). The EmotionPush dialogues are from Facebook Messenger chats by real users which have been anonymized to ensure user privacy¹. For both datasets, dialogue lengths

¹To guarantee high level of anonymization, an automatic Named Entity Recognition (NER) replacement was followed by careful manual inspection and correction

*Contact Author

range from 5 to 24 lines each. A breakdown of the lengths of the dialogues is shown in Table 3.

We employed workers using Amazon Mechanical Turk (aka AMT or MTurk) to annotate the dialogues [Callison-Burch and Dredze, 2010]. Each complete dialogue was offered as a single MTurk Human Intelligence Task (HIT), within which each utterance was read and annotated for emotions by the worker. Each HIT was assigned to five workers. To ensure workers were qualified for the annotation task, we set up a number of requirements: workers had to be from an English-speaking country (Australia, Canada, Great Britain, Ireland, New Zealand, or the US), have a high HIT approval rate (at least 98%), and have already performed a minimum of 2,000 HITs.

Friends	EmotionPush	Both
0.326	0.342	0.345

Table 1: Reliability of Agreement (κ)

In the datasets, each utterance is accompanied by an *annotation* and *emotion*. The annotation contains the raw count of votes for each emotion by the five annotators, with the order of the emotions being Neutral, Joy, Sadness, Fear, Anger, Surprise, Disgust. For example, an annotation of “2000030” denotes that two annotators voted for “neutral”, and three voted for “surprise”.

The labeled emotion is calculated using the absolute majority of votes. Thus, if a specific emotion received three or more votes, then that utterance is labeled with that emotion. If there is no majority vote, the utterance is labeled with “non-neutral” label. In addition to the utterance, annotation, and label, each line in each dialogue includes the speaker’s name (in the case of EmotionPush, a speaker ID was used). The emotion distribution for Friends and EmotionPush, for both training and evaluation data, is shown in Table 2.

We used Fleiss’ kappa measure to assess the reliability of agreement between the annotators [Fleiss, 1971]. The value for κ -statistic is 0.326 and 0.342 for Friends and EmotionPush, respectively. For the combined datasets the value of the κ -statistic is 0.345.²

Sample excerpts from the two datasets, with their annotations and labels, are given in Table 5.

Augmentation NLP tasks require plenty of data. Due to the relatively small number of samples in our datasets, we added more labeled data using a technique developed in [Ostyaikov, 2018] that was used by the winning team in Kaggle’s Toxic Comment Classification Challenge [Lee, 2018]. The augmented datasets are similar to the original data files, but include additional machine-computed utterances for each original utterance. We created the additional utterances using the Google Translate API. Each original utterance was first translated from English into three target languages (German, French, and Italian), and then translated back into English. The resulting utterances were included together in the same

object with the original utterance. These “duplex translations” can sometimes result in the original sentence, but many times variations are generated that convey the same emotions.

Table 4 shows an example utterance (labeled with “Joy”) after augmentation.

3 Challenge Details

A dedicated website for the competition was set up³. The website included instructions, the registration form, schedule, and other relevant details. Following registration, participants were able to download the training datasets.

The label distribution of emotions in our data are highly unbalanced, as can be seen in Figure 1. Due to the small number of three of the labels, participants were instructed to use only four emotions for labels: joy, sadness, anger, and neutral. Evaluation of submissions was done using only utterances with these four labels. Utterances with labels other than the above four (i.e., surprise, disgust, fear or non-neutral) were discarded and not used in the evaluation. Scripts for verifying and evaluating the submissions were made available online⁴. We used micro-F1 as the comparison metric.

4 Submissions

A total of eleven teams submitted their evaluations, and are presented in the online leaderboard. Seven of the teams also submitted technical reports, the highlights of which are summarized below. More details are available in the relevant reports.

IDEA [Huang *et al.*, 2019] Two different BERT models were developed. For Friends, pre-training was done using a sliding window of two utterances to provide dialogue context. Both Next Sentence Prediction (NSP) phase on the complete unlabeled scripts from all 10 seasons of Friends, which are available for download. In addition, the model learned the emotional disposition of each of six main characters in Friends (Rachel, Monica, Phoebe, Joey, Chandler and Ross) by adding a special token to represent the speaker. For EmotionPush, pre-training was performed on Twitter data, as it is similar in nature to chat based dialogues. In both cases, special attention was given to the class imbalance issue by applying “weighted balanced warming” on the loss function.

KU [Yang *et al.*, 2019] BERT is post-trained via Masked Language Model (MLM) and Next Sentence Prediction (NSP) on a corpus consisting of the complete and augmented dialogues of Friends, and the EmotionPush training data. The resulting token embeddings are max-pooled and fed into a dense network for classification. A K -fold cross-validation ensemble with majority voting was used for prediction. To deal with the class imbalance problem, weighted cross entropy was used as a training loss function.

HSU [Luo and Wang, 2019] A pre-trained BERT is fine-tuned using filtered training data which only included the desired labels. Additional augmented data with joy, sadness, and anger labels are also used. BERT is then fed into a

²We used the *irr* package for R.

³<https://sites.google.com/view/emotionx2019>

⁴<http://bit.ly/emotionx2019-verify>

		# utterances	non-neutral	neutral	joy	sadness	anger	disgust	fear	surprise
EmotionPush	Train	14742	1418	9855	2100	514	140	106	42	567
	Test	3536	394	2146	601	110	27	51	18	189
	Subtotal	18278	1812	12001	2701	624	167	157	60	756
Friends	Train	14503	2772	6530	1710	498	759	331	246	1657
	Test	3296	952	1035	505	121	141	58	37	447
	Subtotal	17799	3724	7565	2215	619	900	389	283	2104
Both	Train	29245	4190	16385	3810	1012	899	437	288	2224
	Test	6832	1346	3181	1106	231	168	109	55	636
Total		36077	5536	19566	4916	1243	1067	546	343	2860

Table 2: Emotion Label Distribution




Figure 1: Emotion Label Distribution (Training Datasets)

		Friends		EmotionPush	
Bucket	Lines	Train	Test	Train	Test
1	5-9	250	50	250	60
2	10-14	250	90	250	60
3	15-19	250	70	250	60
4	20-24	250	30	250	60
<i>Dialogues</i>		1000	240	1000	240
<i>Utterances</i>		14503	3296	14742	3536

Table 3: Dialogue Length Distribution and Number of Utterances

standard feed-forward-network with a softmax layer used for classification.

Podlab [Nguyen and Mitchell, 2019] A support vector machine (SVM) was used for classification. Words are ranked using a per-emotion TF-IDF score. Experiments were performed to verify whether the previous utterance would improve classification performance. Input to the Linear SVM was done using one-hot-encoding of top ranking words.

Source	Utterance
Original	"Wow! That's ah, that's pretty nice!"
English \Rightarrow German	"Impressive! That's ah, that's pretty nice!"
English \Rightarrow French	"Wow! It's ah, it's pretty cool!"
English \Rightarrow Italian	"Wow! It's ah, it's cute!"

Table 4: Example of Augmented Utterance

AlexU [Alfons *et al.*, 2019] The classifier uses a pre-trained BERT model followed by a feed-forward neural network with a softmax output. Due to the overwhelming presence of the neutral label, a classifying cascade is employed, where the majority classifier is first used to decide whether the utterance should be classified with "neutral" or not. A second classifier is used to focus on the other emotions (joy, sadness, and anger). Dealing with the imbalanced classes is done through the use of a weighted loss function.

Antenna [Yu *et al.*, 2019] BERT is first used to generate word and sentence embeddings for all utterances. The resulting calculated word embeddings are fed into a Convolutional Neural Network (CNN), and its output is then concatenated with the BERT-generated sentence embeddings. The concatenated vectors are then used to train a bi-directional GRU with

Speaker	Utterance	Annotation	Emotion
Phoebe	“Ohh, that’s too bad.”	2030000	sadness
Ross	“No, I-I’m saying I liked her.”	2010110	non-neutral
Phoebe	“Yeah, y’know what, there are other fish in the sea.”	4010000	neutral
Ross	“Pheebs, I think she’s great. Okay? We’re going out again.”	2100110	non-neutral
Phoebe	“Okay, I hear you! Are you capable of talking about any thing else?”	0000203	disgust
1167038771	“why is that poster so bad”	2000003	disgust
1220662692	“what do you mean its great”	1000040	surprise
1167038771	“oh”	3000020	neutral
1167038771	“Can You Really Trust Anyone”	2010011	non-neutral
1220662692	“lol”	0500000	joy
1167038771	“Can You?”	3001010	neutral
1220662692	“i guess not”	2030000	sadness

Table 5: Dialogue Excerpts from Friends (top) and EmotionPush (bottom)

Rank	Team	Features	F1-micro	Neutral	Joy	Sadness	Anger
1	IDEA	BERT, speaker, transfer, context	81.5	87.3	75.5	59.6	69.8
2	KU	BERT, ensemble, context	79.5	86.0	72.0	51.4	65.6
3	HSU	BERT	79.1	85.4	73.6	55.6	65.0
4	AlexU	BERT, majority classifier	77.0	84.5	72.3	58.0	59.7
5	Antenna	BERT+GRU	75.2	83.5	66.3	49.4	55.9
6	Podlab	SVM, features	74.3	83.8	64.5	42.9	39.5
7	CYUT	bi-GRU	60.3	69.8	52.4	44.3	34.4

Table 6: F-scores for Friends (%)

a residual connection followed by a fully-connected layer, and finally a softmax layer produces predictions. Class imbalance is tackled using focal loss [Lin *et al.*, 2017].

CYUT [Wu and Chien, 2019] A word embedding layer followed by a bi-directional GRU-based RNN. Output from the RNN was fed into a single-node classifier. The augmented dataset was used for training the model, but “neutral”-labeled utterances were filtered to deal with class imbalance.

5 Results

The submissions and the final results are summarized in Tables 6 and 7. Two of the submissions did not follow up with technical papers and thus they do not appear in this summary. We note that the top-performing models used BERT, reflecting the recent state-of-the-art performance of this model in many NLP tasks. For Friends and EmotionPush the top micro-F1 scores were 81.5% and 88.5% respectively.

6 Evaluation & Discussion

An evaluation summary of the submissions is available in Tables 6 and 7. We only present the teams that submitted technical reports. A full leaderboard that includes all the teams is available on the challenge website. This section highlights some observations related to the challenge. Identical utterances can convey different emotions in different contexts. A few of the models incorporated the dialogue context into the model, such as the models proposed by teams **IDEA** and **KU**.

Deep Learning Models. Most of the submissions used deep learning models. Five of the models were based on the BERT architecture, with some using pre-trained BERT. Some of the submissions enhanced the model by adding context and speaker related encoding to improve performance. We also received submissions using more traditional networks such as CNN, as well as machine learning classics such as SVM. The results demonstrate that domain knowledge, feature engineering, and careful application of existing methodologies is still paramount for building successful machine learning models.

Unbalanced Labels. Emotion detection in text often suffers from a data imbalance problem, our datasets included. The teams used two approaches to deal with this issue. Some used a class-balanced loss functions while others under-sampled classes with majority label “neutral”. Classification performance of underrepresented emotions, especially sadness and anger, is low compared to the others. This is still a challenge, especially as some real-world applications are dependent on detection of specific emotions such as anger and sadness.

Emotional Model and Annotation Challenges. The discrete 6-emotion model and similar models are often used in emotion detection tasks. However, such 1-out-of-n models are limited in a few ways: first, expressed emotions are often not discrete but mixed (for example, surprise and joy or surprise and anger are often manifested in the same utterance). This leads to more inter-annotator disagreement, as annotators can only select one emotion. Second, there are additional

Rank	Team	Approach	F1-micro	Neutral	Joy	Sadness	Anger
1	IDEA	BERT, speaker, transfer, context	88.5	92.7	77.7	60.4	47.4
2	KU	BERT, context, ensemble	87.6	92.2	73.9	61.8	40.0
3	HSU	BERT	86.2	91.3	73.8	62.7	39.1
4	Podlab	SVM, features	85.7	91.3	67.8	52.0	22.9
5	AlexU	BERT, majority classifier	84.1	90.1	75.8	52.4	30.8
6	CYUT	bi-GRU	71.5	81.1	58.8	28.6	10.5
7	Antenna	BERT+GRU	70.2	82.2	27.9	17.9	6.7

Table 7: F-scores for EmotionPush (%)

	Friends	EmotionPush	Both
Joy	0.401	0.454	0.429
Neutral	0.350	0.349	0.373
Surprise	0.330	0.266	0.322
Sadness	0.291	0.324	0.307
Anger	0.319	0.212	0.306
Disgust	0.184	0.160	0.179
Fear	0.184	0.132	0.176

Table 8: Per-emotion Reliability of Agreement (κ)

emotional states that are not covered by the basic six emotions but are often conveyed in speech and physical expressions, such as desire, embarrassment, relief, and sympathy. This is reflected in feedback we received from one of the AMT workers: “*I am doing my best on your HITS. However, the emotions given (7 of them) are a lot of times not the emotion I’m reading (such as questioning, happy, excited, etc). Your emotions do not fit them all...*”.

To further investigate, we calculated the per-emotion κ -statistic for our datasets in Table 8. We see that for some emotions, such as disgust and fear (and anger for EmotionPush), the κ -statistic is poor, indicating ambiguity in annotation and thus an opportunity for future improvement. We also note that there is an interplay between the emotion label distribution, per-emotion classification performance, and their corresponding κ scores, which calls for further investigation.

Data Sources. One of the main requirements of successful training of deep learning models is the availability of high-quality labeled data. Using AMT to label data has proved to be useful. However, current data is limited in quantity. In addition, more work needs to be done in order to measure, evaluate, and guarantee annotation quality. In addition, the Friends data is based on an American TV series which emphasizes certain emotions, and it remains to be seen how to transfer learning of emotions to other domains.

7 Acknowledgment

This research is partially supported by Ministry of Science and Technology, Taiwan, under Grant no. MOST108-2634-F-001-004- and MOST107-2218-E-002-009-.

References

- [Alfons *et al.*, 2019] Meena Alfons, Marwan Torki, and Nagwa El-Makky. EmotionX-AlexU: Cascade BERT-based emotion classifier. Technical report, 2019.
- [Bazzanella, 2004] Carla Bazzanella. Emotions, language, and context. In *Emotion in Dialogic Interaction*, page 59. John Benjamins Publishing Company, 2004.
- [Callison-Burch and Dredze, 2010] Chris Callison-Burch and Mark Dredze. Creating speech and language data with amazon’s mechanical turk. In *Proceedings of the NAACL HLT 2010 Workshop on Creating Speech and Language Data with Amazon’s Mechanical Turk*, pages 1–12. Association for Computational Linguistics, 2010.
- [Chen *et al.*, 2018] Sheng-Yeh Chen, Chao-Chun Hsu, Chuan-Chun Kuo, Lun-Wei Ku, et al. Emotionlines: An emotion corpus of multi-party conversations. *arXiv preprint arXiv:1802.08379*, 2018.
- [Ekman, 1999] Paul Ekman. Basic emotions. *Handbook of cognition and emotion*, pages 45–60, 1999.
- [Fleiss, 1971] Joseph L Fleiss. Measuring nominal scale agreement among many raters. *Psychological bulletin*, 76(5):378, 1971.
- [Hsu and Ku, 2018] Chao-Chun Hsu and Lun-Wei Ku. SocialNLP 2018 EmotionX challenge overview: Recognizing emotions in dialogues. In *Proceedings of the Sixth International Workshop on Natural Language Processing for Social Media*, pages 27–31, 2018.
- [Huang *et al.*, 2019] Yen Hao Huang, Ssu-Rui Lee, Mau-Yun Ma, Yi Hsin Chen, Ya-Wen Yu, and Yi-Shin Chen. EmotionX-IDEA: Emotion BERT—an affectional model for conversation. Technical report, 2019.
- [Lee, 2018] Chun Ming Lee. 1st place solution overview. <https://www.kaggle.com/c/jigsaw-toxic-comment-classification-challenge/discussion/52557>, Mar 2018.
- [Lin *et al.*, 2017] Tsung-Yi Lin, Priya Goyal, Ross Girshick, Kaiming He, and Piotr Dollár. Focal loss for dense object detection. In *Proceedings of the IEEE international conference on computer vision*, pages 2980–2988, 2017.
- [Luo and Wang, 2019] Linkai Luo and Yue Wang. EmotionX-HSU: Adopting pre-trained BERT for emotion classification. Technical report, 2019.

- [Nguyen and Mitchell, 2019] Andrew Nguyen and Lewis Mitchell. EmotionX-Podlab: A linear SVM for emotion classification. Technical report, 2019.
- [Ostyakov, 2018] Pavel Ostyakov. A simple technique for extending dataset. <https://www.kaggle.com/c/jigsaw-toxic-comment-classification-challenge/discussion/48038>, Jan 2018.
- [Seyeditabari *et al.*, 2018] Armin Seyeditabari, Narges Tabari, and Wlodek Zadrozny. Emotion detection in text: a review. *arXiv preprint arXiv:1806.00674*, 2018.
- [Wu and Chien, 2019] Shih-Hung Wu and Sheng-Lun Chien. EmotionX-CYUT: Attention RNN model with least context for emotion detection in dialogues. Technical report, 2019.
- [Yang *et al.*, 2019] Kisu Yang, Dongyub Lee, Taesun Whang Seolhwa Lee, and Heuiseok Lim. EmotionX-KU: BERT-Max based contextual emotion classifier. Technical report, 2019.
- [Yu *et al.*, 2019] Zhuoran Yu, Yuhong Wang, Zhanhao Liu, and Xiang Cheng. EmotionX-Antenna: An emotion detector with residual GRU and text CNN. Technical report, 2019.