

Apache Storm and Spark Streaming Compared

P. Taylor Goetz, Hortonworks
@ptgoetz

Honestly...

- I know a lot more about Apache Storm than I do Apache Spark Streaming.
- I've been involved with Apache Storm, in one way or another, since it was open-sourced.
- I'm admittedly biased.

But...

- A number of articles/papers comparing Apache Storm and Spark Streaming are inaccurate in terms of Storm's features and performance characteristics.
- Code and configuration for those studies is not available, so independent verification is impossible.
- Claims don't match real-world observations.

But...

- There is an inherent “Home Team Advantage” in any benchmark comparison.
- Without open source code, any benchmark claims are essentially marketing fluff, and should be taken with a grain or two of NaCl.
- Any benchmark claim should be independently verifiable.

Spark Streaming Paper

- Compares Spark Streaming (Micro-Batch) to Core Storm (One-at-a-Time)
- A more appropriate comparison would have been with Storm's Trident (Micro-Batch) API
- Trident mentioned only in passing (on pages 3 and 12)

Spark Streaming Paper

- Benchmark code/configuration not publicly available
- Performance claims not independently verifiable

Spark Streaming Paper

- Granted, the Spark Streaming paper is almost 2 years old and written at a time when Trident was relatively new.
- However, that paper is often cited when comparing Apache Storm and Spark Streaming, particularly in terms of performance.
- A lot can change in 2 years.

Streaming and batch
processing are
fundamentally different.

Batch vs. Streaming

- ***Storm*** is a stream processing framework that also does micro-batching (Trident).
- ***Spark*** is a batch processing framework that also does micro-batching (Spark Streaming).

Batch vs. Streaming



Batch vs. Streaming



Apache Storm: Two Streaming APIs

Core Storm (Spouts and Bolts)

- One at a Time
- Lower Latency
- Operates on Tuple Streams

Trident (Streams and Operations)

- Micro-Batch
- Higher Throughput
- Operates on Streams of Tuple Batches and Partitions

Language Options

| Core Storm | Storm Trident | Spark Streaming |
|--|--|---|
| <ul style="list-style-type: none">• Java• Clojure• Scala• Python• Ruby• others* | <ul style="list-style-type: none">• Java• Clojure• Scala | <ul style="list-style-type: none">• Java• Scala• Python |

*Storm's Multi-Lang feature allows the use of virtually any programming language.

Reliability Models

| | Core Storm | Storm Trident | Spark Streaming |
|---------------|------------|---------------|-----------------|
| At Most Once | Yes | Yes | No |
| At Least Once | Yes | Yes | No* |
| Exactly Once | No | Yes | Yes* |

*In some node failure scenarios, Spark Streaming falls back to at-least-once processing or data loss.

Programing Model

| | Core Storm | Storm Trident | Spark Streaming |
|--------------------------------|-----------------------|---|-----------------------------------|
| Stream Primitive | Tuple | Tuple, Tuple Batch, Partition | DStream |
| Stream Source | Spouts | Spouts, Trident Spouts | HDFS, Network |
| Computation/ Transformation | Bolts | Filters, Functions, Aggregations, Joins | Transformation, Window Operations |
| Stateful Operations | No (roll your own) | Yes | Yes |
| Output/ Persistence | Bolts | State, MapState | foreachRDD |

Production Deployments

Apache Storm

- Too many to list

[http://
storm.incubator.apache.org/
documentation/Powered-
By.html](http://storm.incubator.apache.org/documentation/Powered-By.html)

Spark Streaming

- Sharethrough

[http://
engineering.sharethrough.com/blog/
2014/06/27/sharethrough-at-spark-
summit-2014-spark-streaming-for-
realtime-auctions/](http://engineering.sharethrough.com/blog/2014/06/27/sharethrough-at-spark-summit-2014-spark-streaming-for-realtime-auctions/)

Support

| | Apache Storm | Spark | Spark Streaming |
|--------------------------|-------------------|---------------------------------------|-----------------------------|
| Hadoop Distro | Hortonworks, MapR | Cloudera, MapR, Hortonworks (preview) | Hortonworks, Cloudera, MapR |
| Resource Management | YARN, Mesos | YARN, Mesos | YARN*, Mesos |
| Provisioning/ Monitoring | Apache Ambari | Cloudera Manager | ? |

*With issues: <http://spark-summit.org/wp-content/uploads/2014/07/Productionizing-a-247-Spark-Streaming-Service-on-YARN-Ooyala.pdf>

Failure Scenarios

Worker Failure: Spark Streaming

"So if a worker node fails, then the system can recompute the lost from the the left over copy of the input data. However, if the worker node where a network receiver was running fails, **then a tiny bit of data may be lost**, that is, the data received by the system but not yet replicated to other node(s)."

Only HDFS-backed data sources are fully fault tolerant.

<https://spark.apache.org/docs/latest/streaming-programming-guide.html#fault-tolerance-properties>

Worker Failure: Spark Streaming

Solution?: Write Ahead Logs (SPARK-3129)

- Enabling WAL requires DFS (HDFS, S3) — no such requirement with Storm
- Incurs a performance penalty that adds to overall latency
- ***Full fault tolerance still requires a data source that can replay data (e.g. Kafka)***
- Architectural band aid?

<https://databricks.com/blog/2015/01/15/improved-driver-fault-tolerance-and-zero-data-loss-in-spark-streaming.html>