

Building microservices with Scala, functional domain models and Spring Boot

Chris Richardson

Author of POJOs in Action

Founder of the original CloudFoundry.com

🐦 @crichardson

chris@chrisrichardson.net



<http://plainoldobjects.com>



Presentation goal

Share my experiences with building an application using Scala, functional domain models, microservices, event sourcing, CQRS, and Spring Boot

About Chris



About Chris

- Founder of a buzzword compliant (stealthy, social, mobile, big data, machine learning, ...) startup
- Consultant helping organizations improve how they architect and deploy applications using cloud, micro services, polyglot applications, NoSQL, ...

Agenda

- Why build event-driven microservices?
- Overview of event sourcing
- Designing microservices with event sourcing
- Implementing queries in an event sourced application
- Building and deploying microservices

Let's imagine that you are
building a banking app...

Domain model

Account

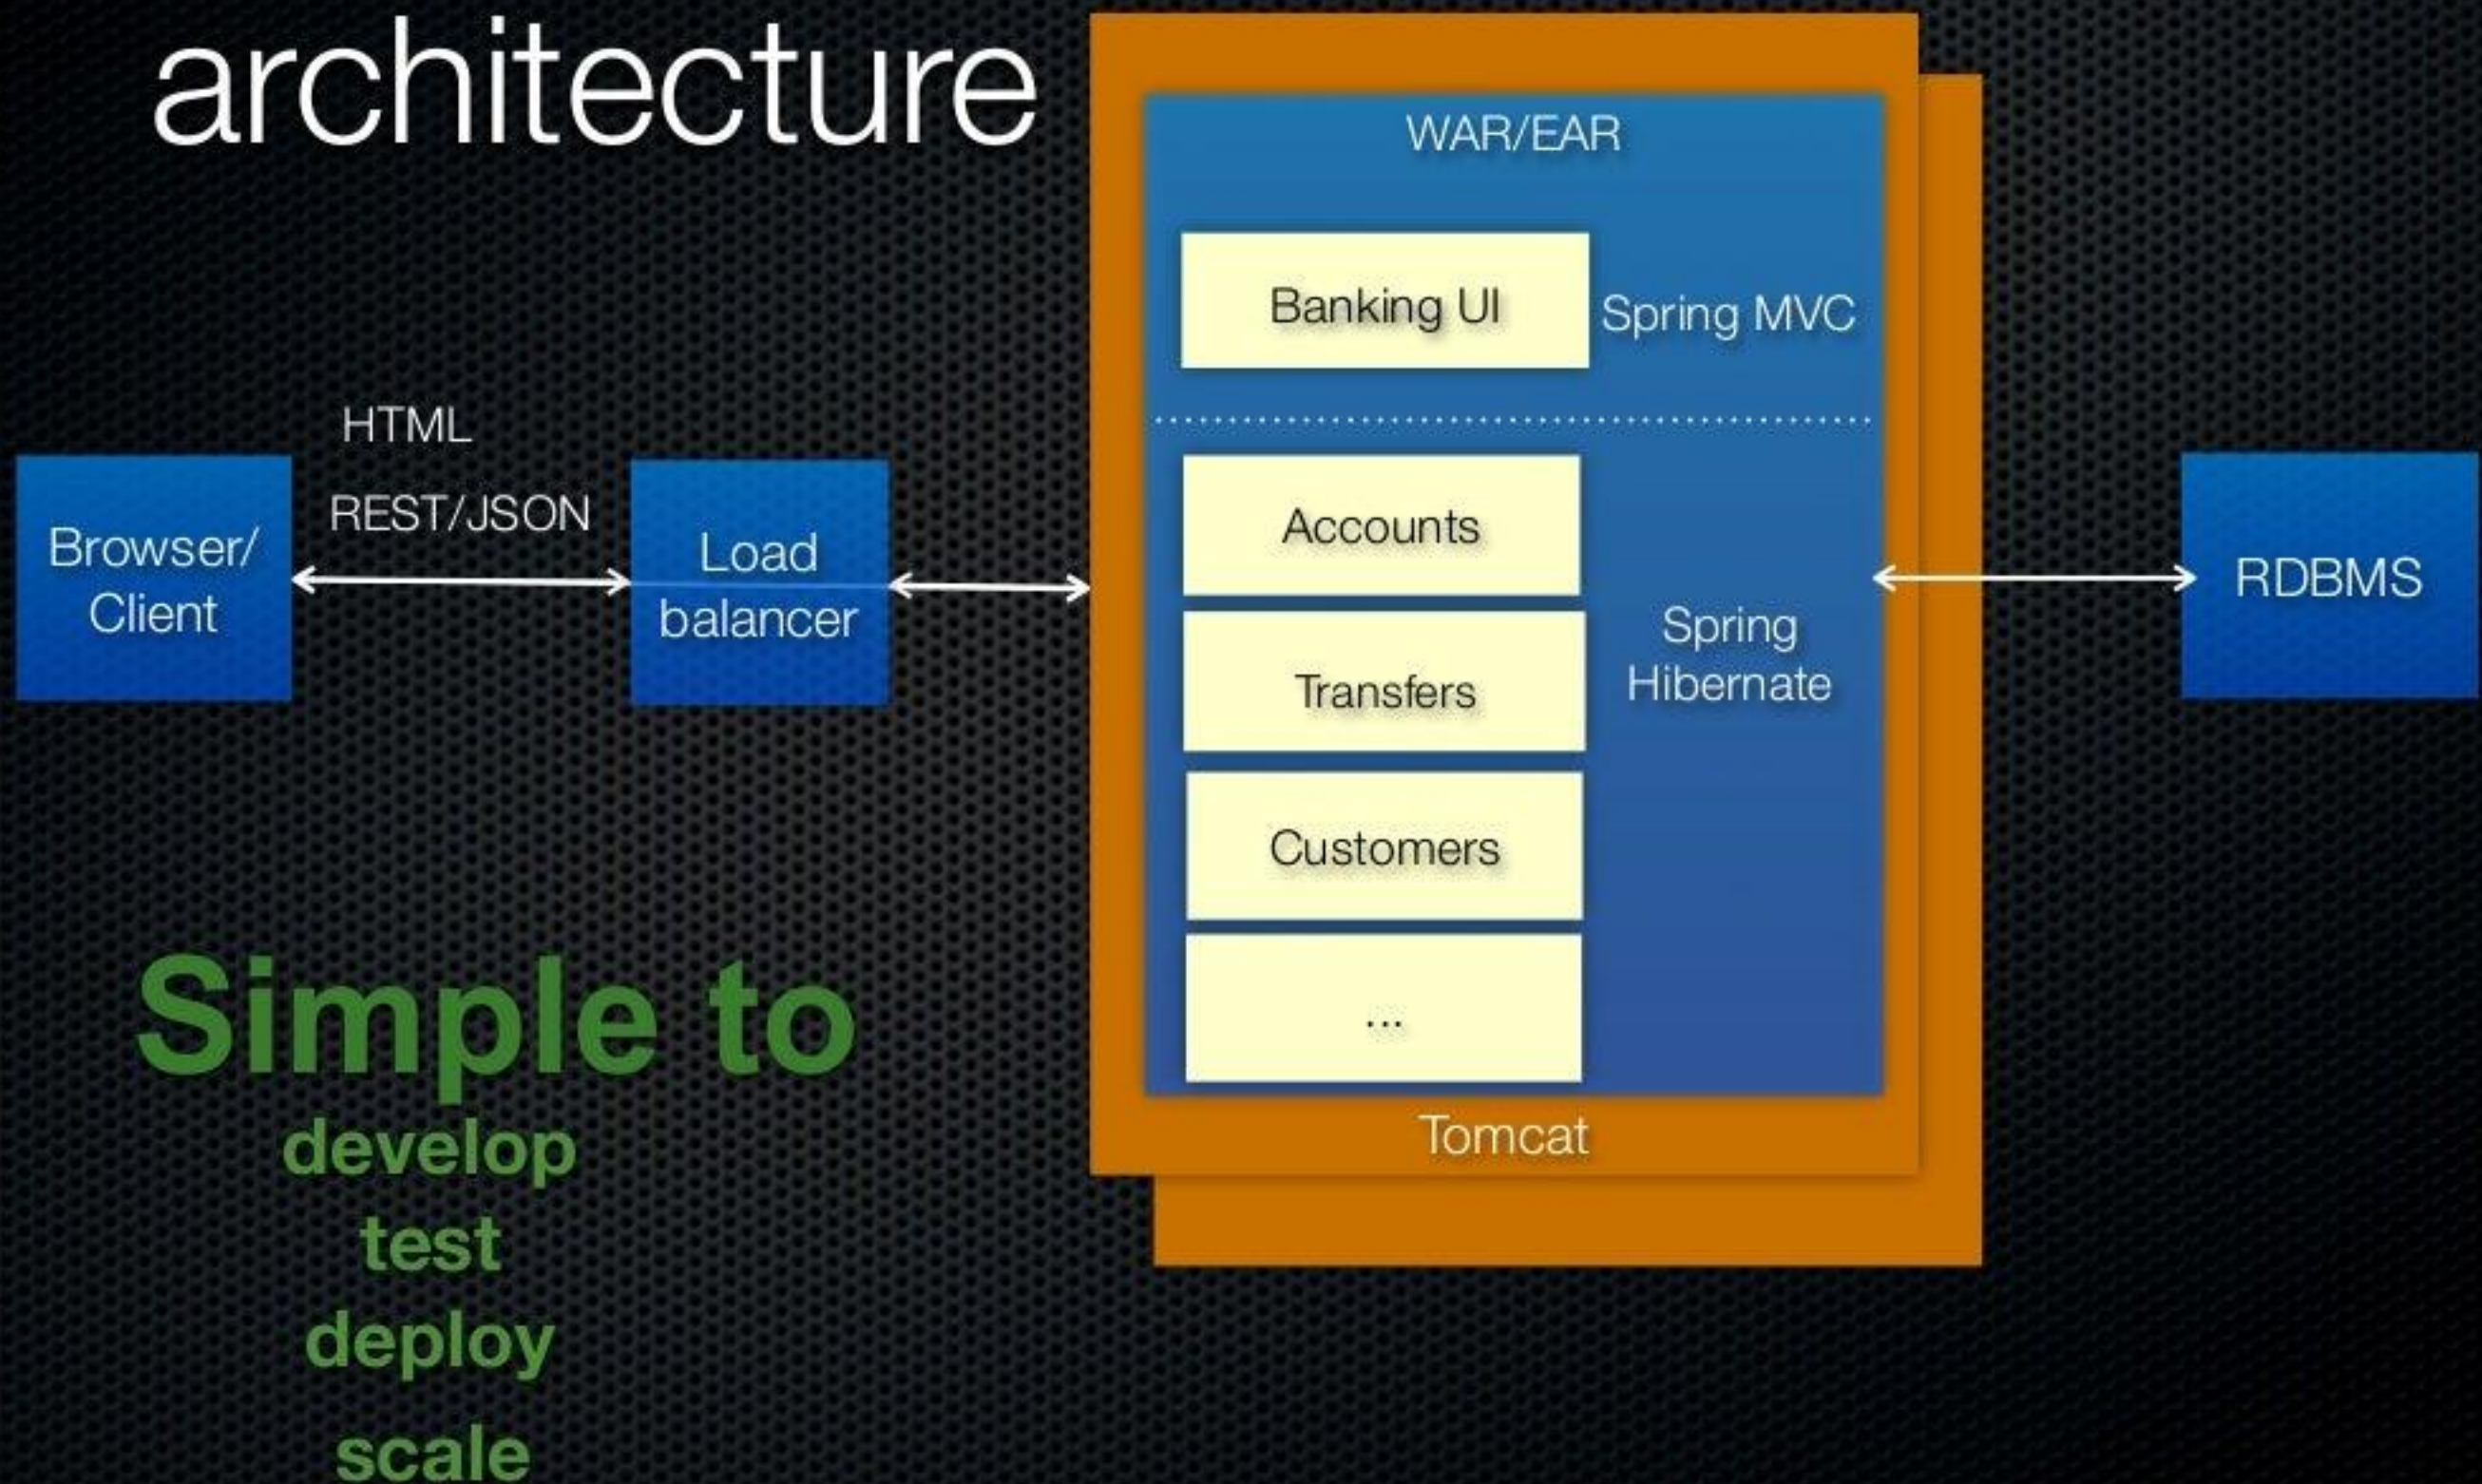
balance

open(initialBalance)
debit(amount)
credit(amount)

MoneyTransfer

fromAccountId
toAccountId
amount

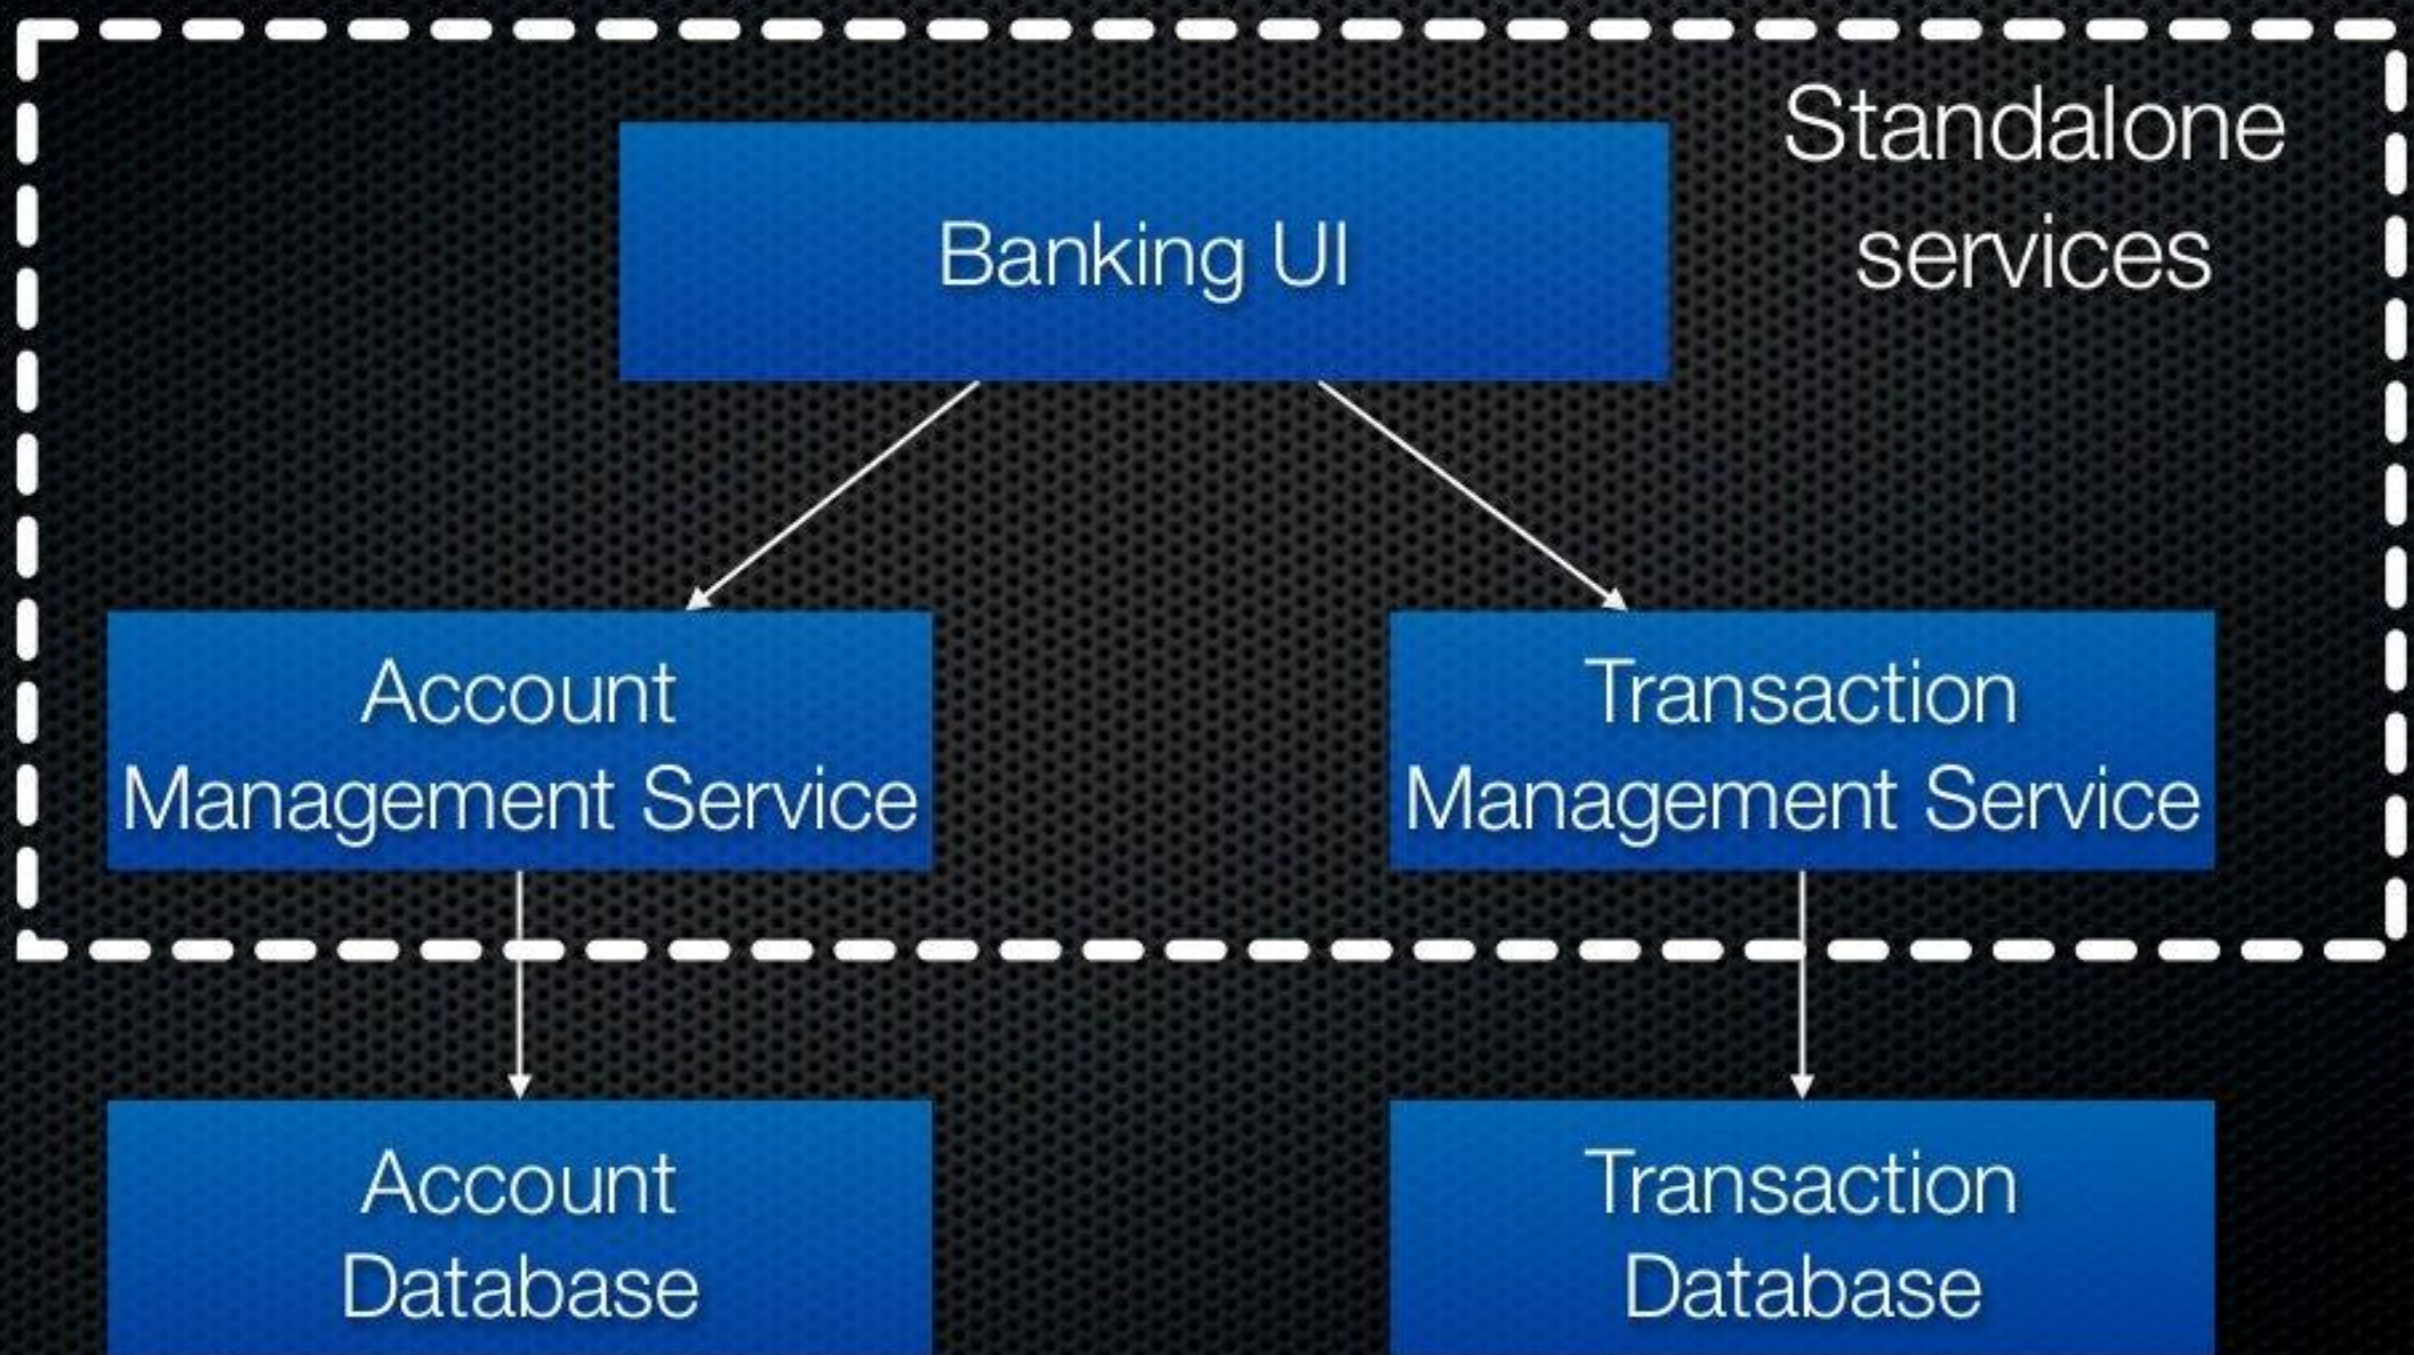
Traditional application architecture



Problem #1: monolithic architecture

- Intimidates developers
- Obstacle to frequent deployments
- Overloads your IDE and container
- Obstacle to scaling development
- Requires long-term commitment to a technology stack

Solution #1: use a microservice architecture



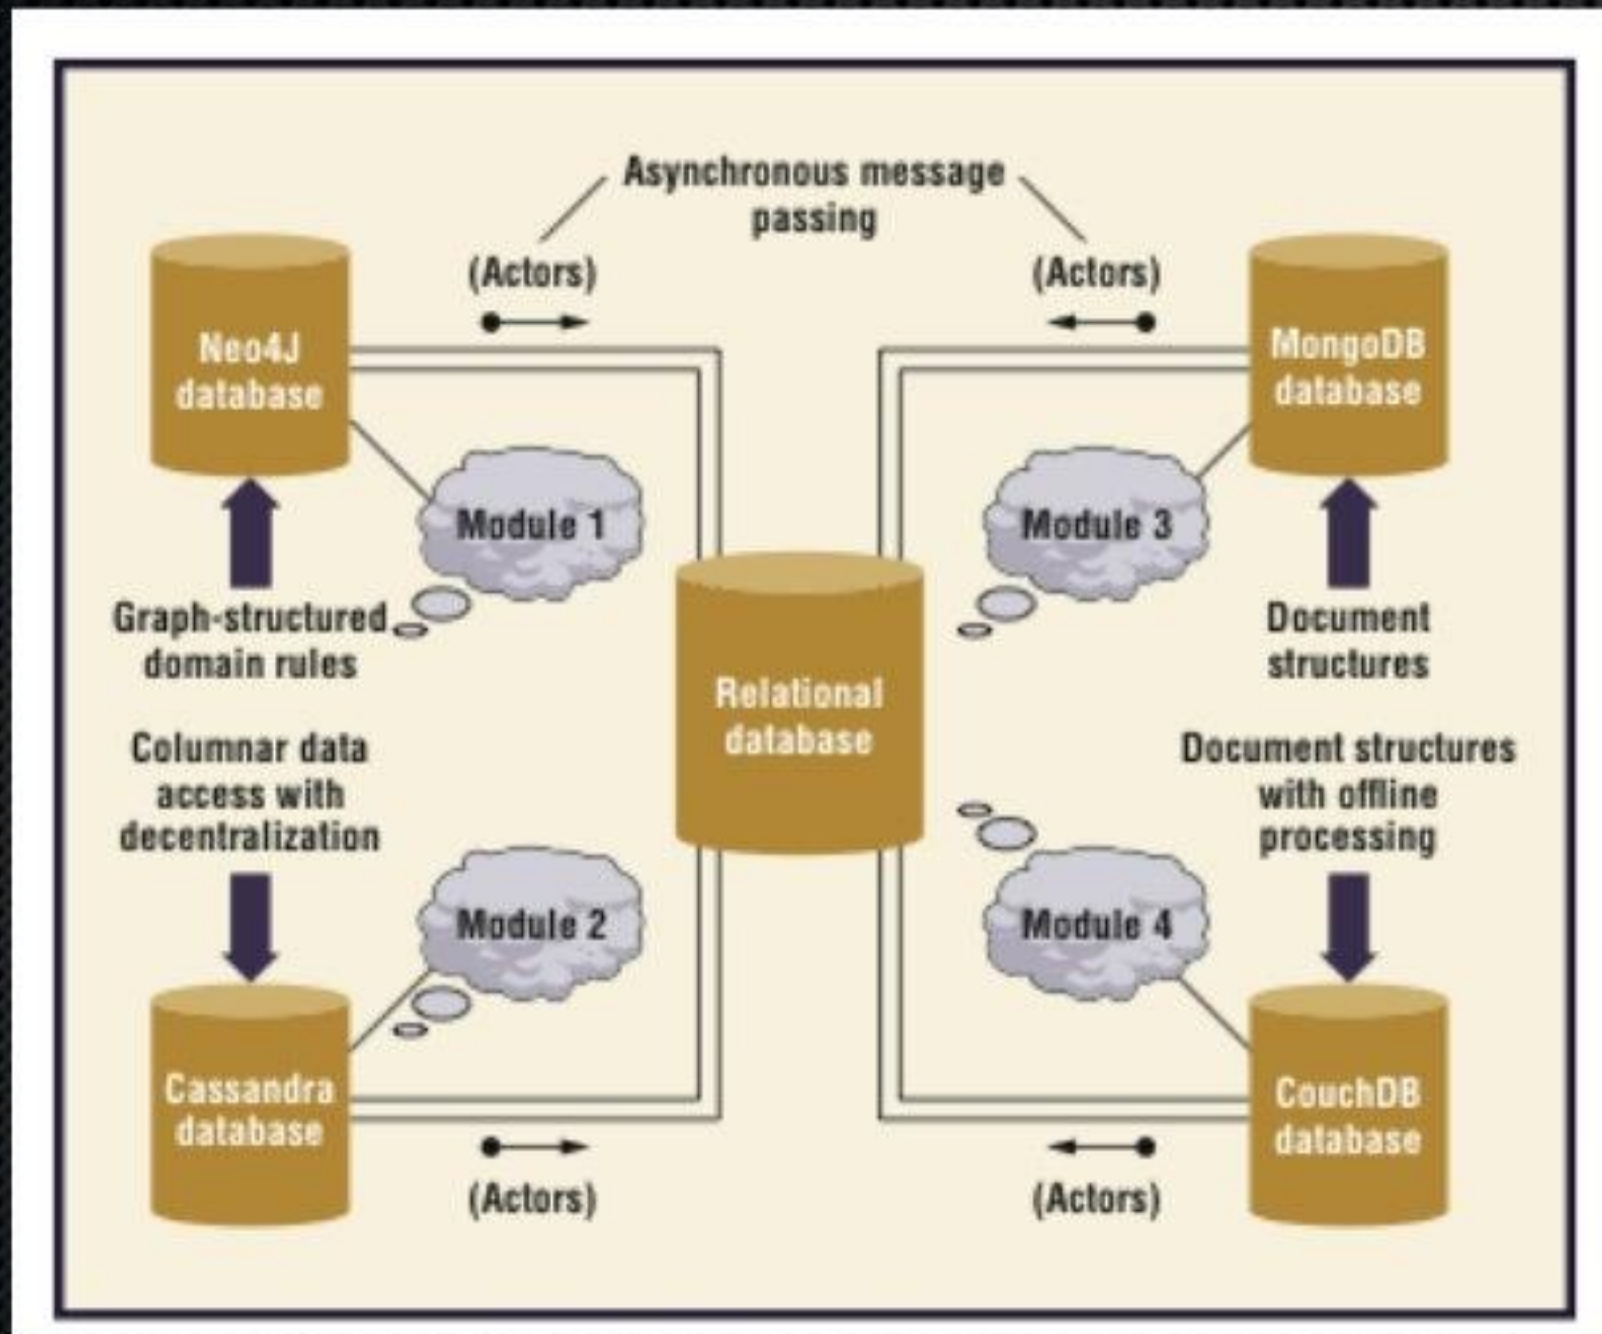
Problem #2: relational databases

- Scalability
- Distribution
- Schema updates
- O/R impedance mismatch
- Handling semi-structured data

Solution #2: use NoSQL databases

- Avoids the limitations of RDBMS
- For example,
 - text search \Rightarrow Solr/Cloud Search
 - social (graph) data \Rightarrow Neo4J
 - highly distributed/available database \Rightarrow Cassandra
 - ...

Different modules use different databases



But now we have problems
with data consistency!

Problem #3: Microservices = distributed data management

- Each microservice has it's own database
 - Some data is replicated and must be kept in sync
 - Business transactions must update data owned by multiple services
- Tricky to implement reliably without 2PC

Problem #4: NoSQL = ACID-free, denormalized databases

- Limited transactions, i.e. no ACID transactions
 - Tricky to implement business transactions that update multiple rows, e.g. <http://bit.ly/mongo2pc>
- Limited querying capabilities
 - Requires denormalized/materialized views that must be synchronized
 - Multiple datastores (e.g. DynamoDB + Cloud Search) that need to be kept in sync

Solution to #3/#4: Event-based architecture to the rescue

- Microservices publish events when state changes
- Microservices subscribe to events
 - Synchronize replicated data
 - Maintains eventual consistency across multiple aggregates (in multiple datastores)

Eventually consistent money transfer

transferMoney()



MoneyTransferService

AccountService

Subscribes to:

AccountDebitedEvent
AccountCreditedEvent

publishes:

MoneyTransferCreatedEvent
DebitRecordedEvent

Subscribes to:

MoneyTransferCreatedEvent
DebitRecordedEvent

Publishes:

AccountDebitedEvent
AccountCreditedEvent

Message Bus

To maintain consistency a service
must
atomically publish an event
whenever
a domain object changes

How to generate events reliably?

- Database triggers, Hibernate event listener, ...
 - Reliable BUT
 - Not with NoSQL
 - Disconnected from the business level event
 - Limited applicability
- Ad hoc event publishing code mixed into business logic
 - Messy code, poor separation of concerns
 - Unreliable, e.g. too easy to forget to publish an event