

Chicago Apache Flink Meetup

June 30th, 2015

Overview of **Apache Flink**: Next-Gen Big Data Analytics Framework

By Slim Baltagi @SlimBaltagi



Apache Flink

Agenda

- I. What is Apache **Flink stack** and how it **fits** into the **Big Data ecosystem**?
- II. How Apache **Flink integrates** with Apache **Hadoop** and other open source tools?
- III. Why Apache **Flink** is an **alternative** to Apache **Hadoop MapReduce**, Apache **Storm** and Apache **Spark**?
- IV. **Who is using** Apache **Flink**?
- V. Where to **learn more** about Apache **Flink**?
- VI. What are some Key **Takeaways**?

I. What is Apache Flink stack and how it fits into the Big Data ecosystem?

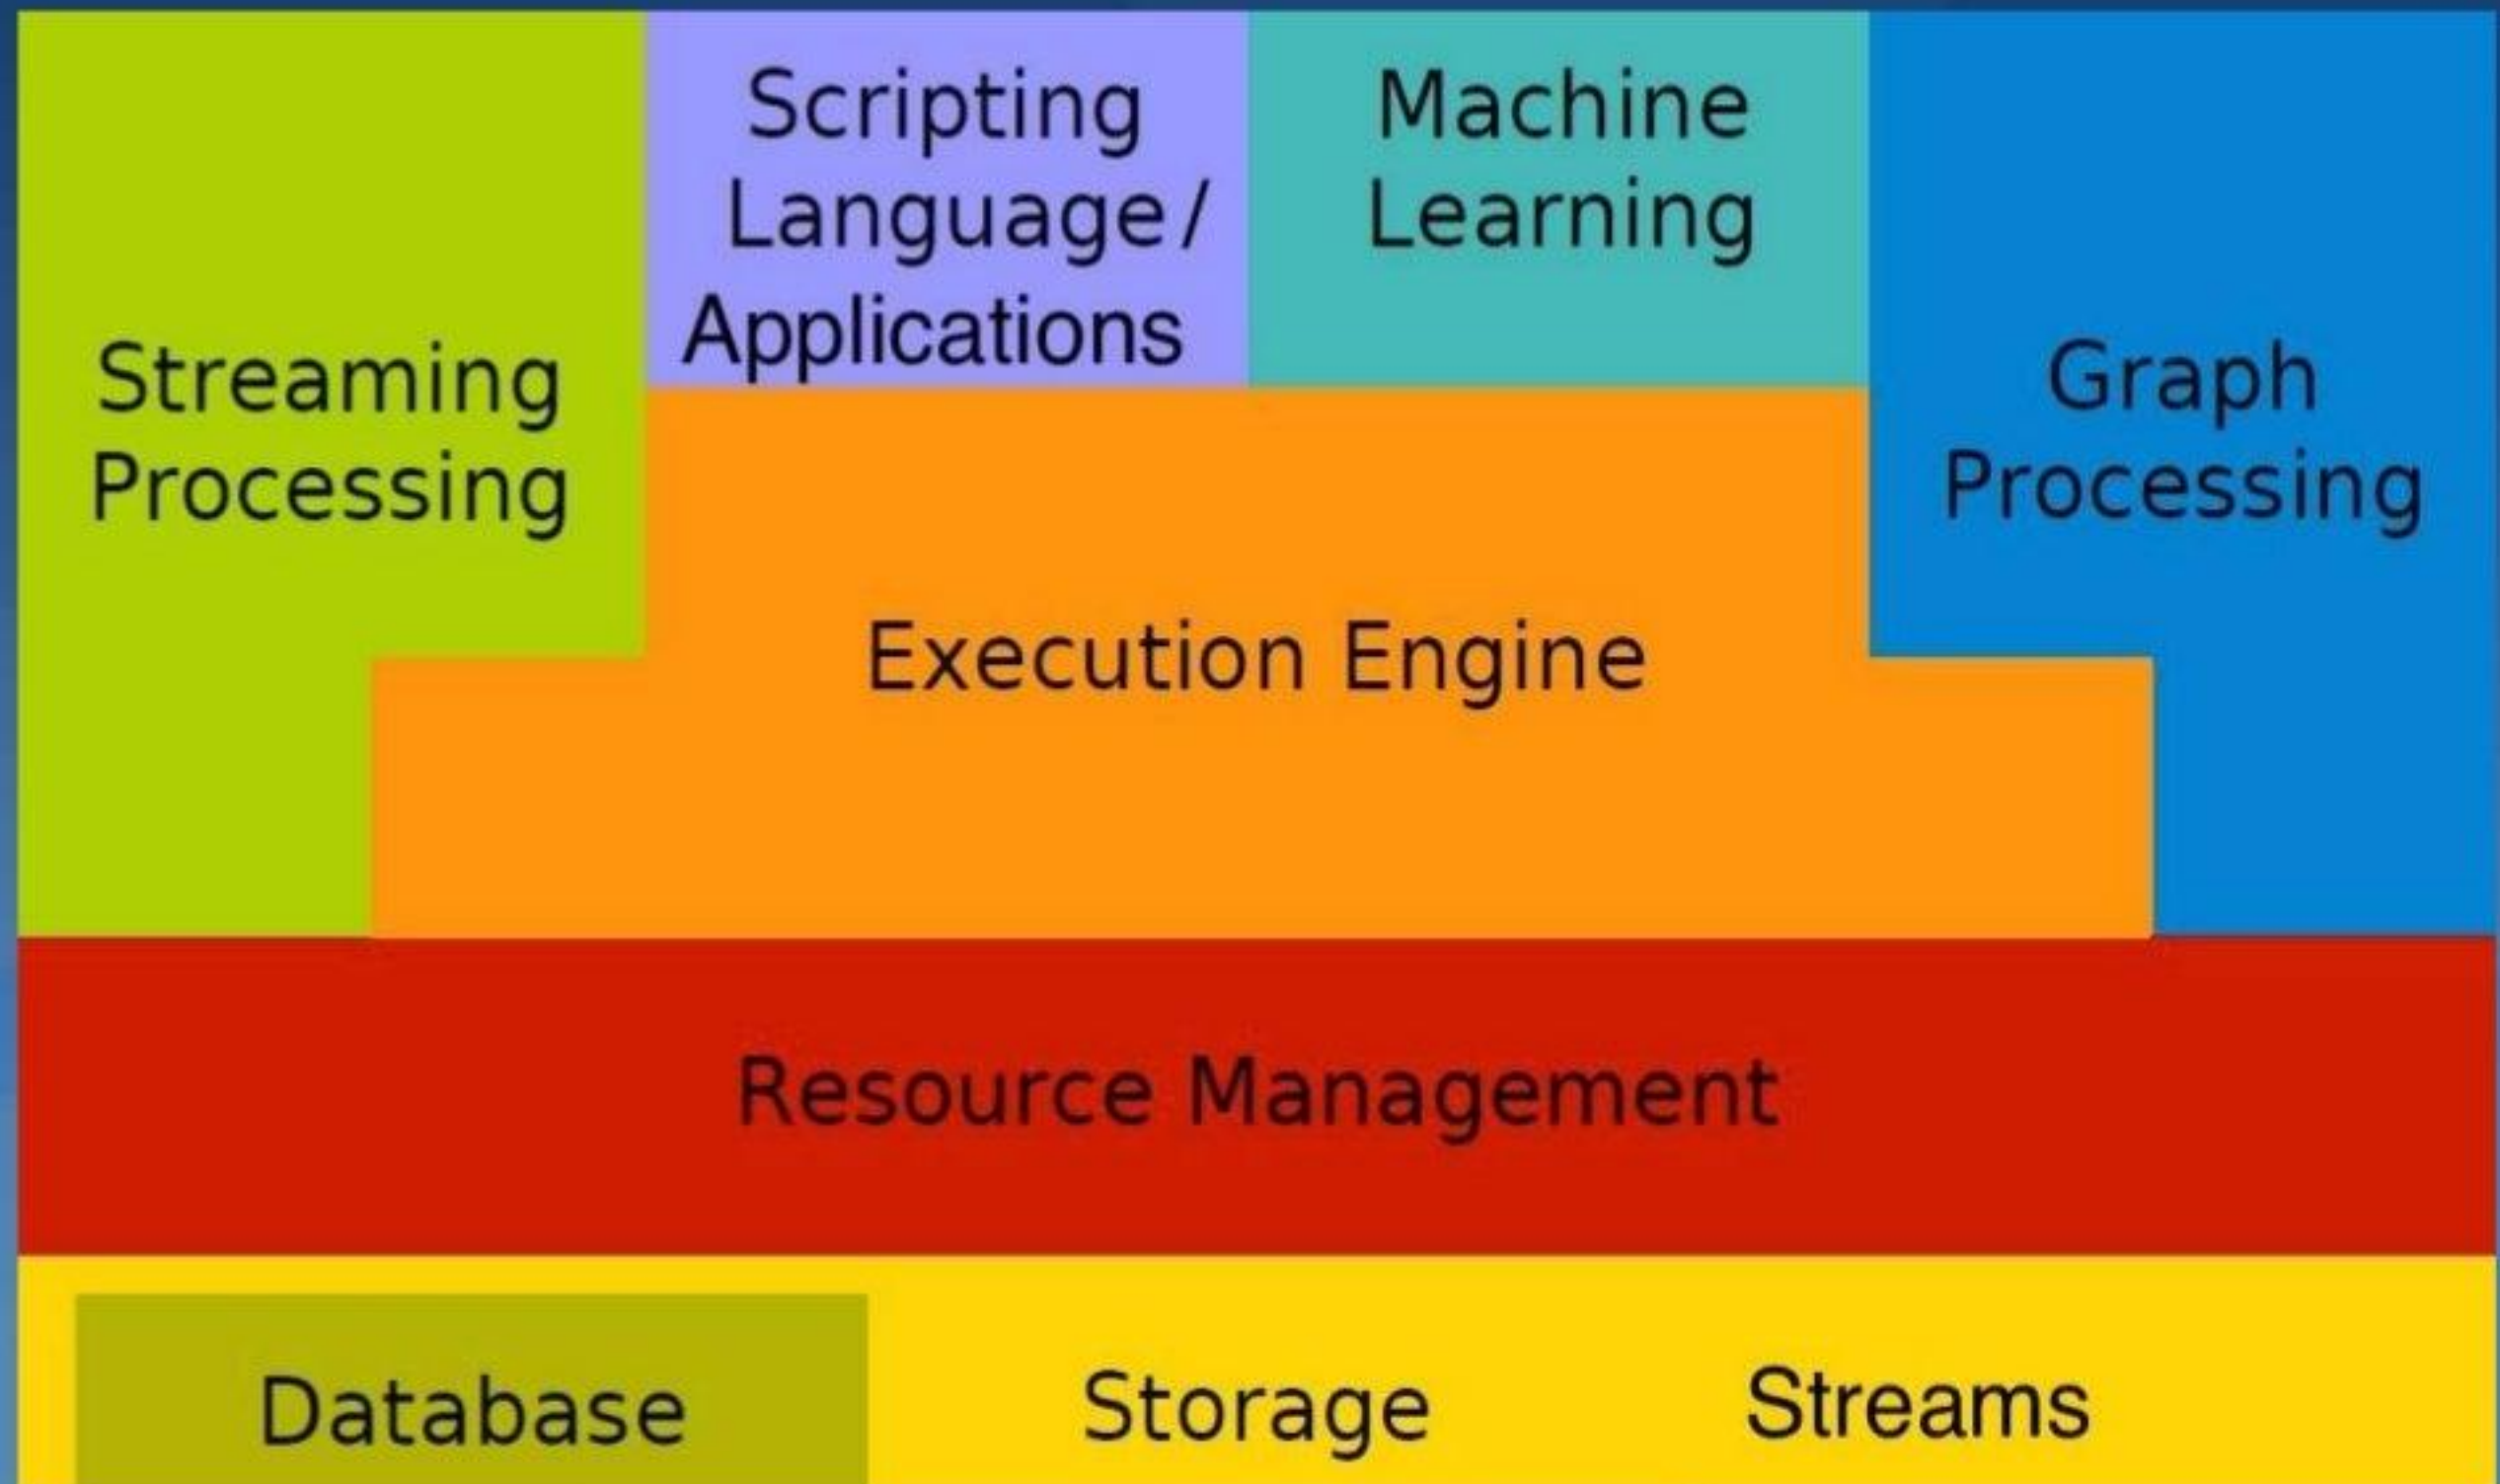
1. What is **Big Data**?
2. What is a typical **Big Data Analytics Stack**?
3. What is **Apache Flink**?
4. What is **Flink Execution Engine**?
5. What are **Flink APIs**?
6. What are **Flink Domain Specific Libraries**?
7. How Flink offers **Interactive Data Analysis**?
8. What is **Flink Architecture**?
9. What is **Flink Programming Model**?

1. What is Big Data?

“Big Data refers to **data sets** large enough and **data streams** fast enough, from **heterogeneous data sources**, that has **outpaced our capability** to store, process, analyze, and understand.”



2. What is a typical Big Data Analytics Stack: Hadoop, Spark, Flink, ...?



3. What is Apache Flink?

- **Apache Flink**, like Apache Hadoop and Apache Spark, is a community-driven open source framework for **distributed Big Data Analytics**. Apache Flink engine exploits **data streaming** and **in-memory processing** and **iteration operators** to improve performance.
- **Apache Flink** has its origins in a research project called **Stratosphere** of which the idea was conceived in 2008 by professor **Volker Markl** from the **Technische Universität Berlin** in **Germany**.
- In German, Flink means **agile** or **swift**. Flink joined the Apache incubator in April 2014 and graduated as an **Apache Top Level Project (TLP)** in **December 2014**.



3. What is Apache Flink?



The Apache **Flink** framework, **written in Java**, provides:

1. Big data processing engine: **distributed and scalable streaming dataflow engine**
2. Several **APIs** in Java/Scala/Python:
 - DataSet API – **Batch** processing
 - DataStream API – **Real-Time streaming** analytics
 - Table API - **Relational** Queries
3. **Domain-Specific Libraries:**
 - FlinkML: **Machine Learning** Library for Flink
 - Gelly: **Graph** Library for Flink
4. **Shell for interactive data analysis**

Key Vision of Apache Flink

**Draws on concepts
from
MPP Database
Technology**

- **Declarativity**
- **Query optimization**
- **Efficient parallel in-memory and out-of-core algorithms**

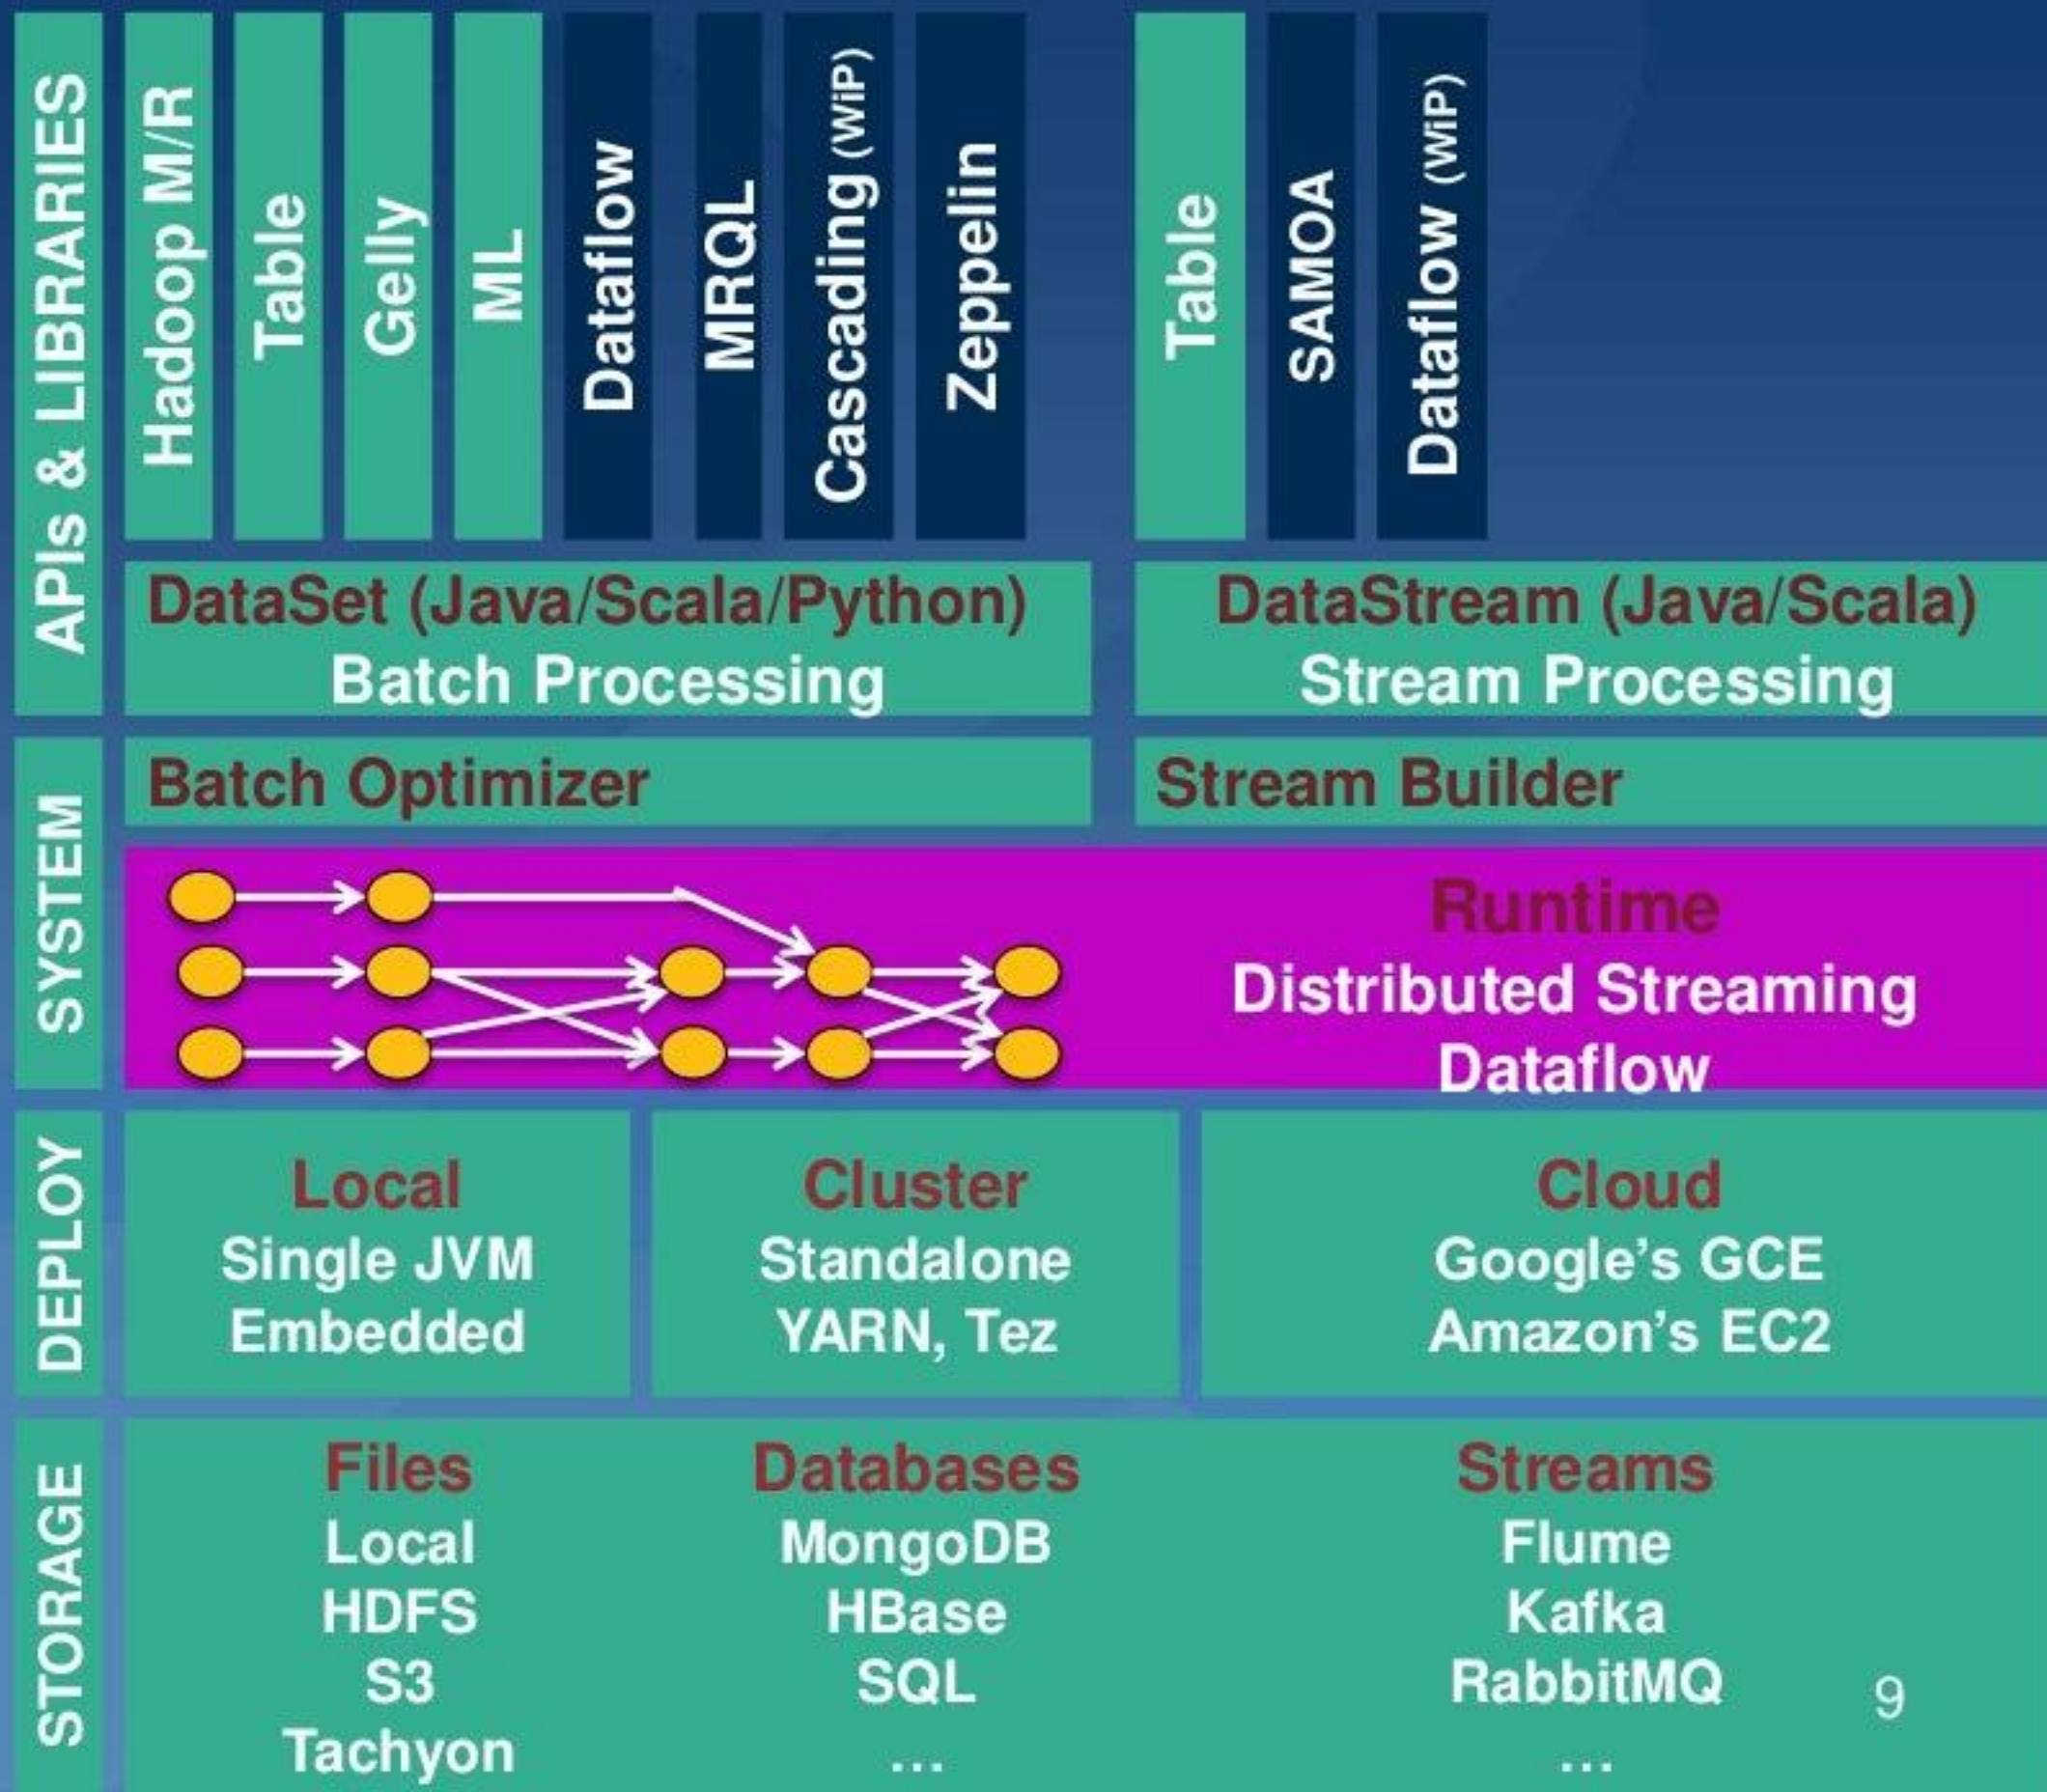
Add

- **Streaming**
- **Iterations**
- **Advanced Dataflows**
- **General APIs**

**Draws on concepts
from
Hadoop MapReduce
Technology**

- **Massive scale-out**
- **User Defined Functions**
- **Complex data types**
- **Schema on read**

What is Apache Flink stack?

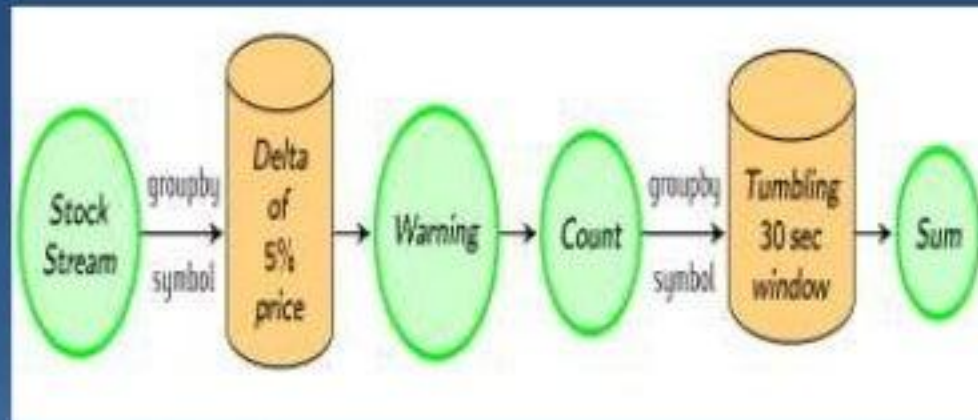


4. What is Flink Execution Engine?

The core of Flink is a **distributed and scalable streaming dataflow engine** with some **unique features**:

1. **True streaming** capabilities: Execute everything as streams
2. **Native iterative** execution: Allow some cyclic dataflows
3. **Handling of mutable state**
4. **Custom memory manager**: Operate on managed memory
5. **Cost-Based Optimizer**: for both batch and stream processing

The **only hybrid** (Real-Time Streaming + Batch) open source distributed data processing **engine** supporting **many use cases**:



Real-Time stream processing



User	Item			
	W	X	Y	Z
A		4.5	2.0	
B	4.0		3.5	
C		5.0		2.0
D		3.5	4.0	1.0

Rating Matrix

=

User	Item			
	W	X	Y	Z
A	1.2	0.8		
B	1.4	0.9		
C	1.5	1.0		
D	1.2	0.8		


User Matrix

X


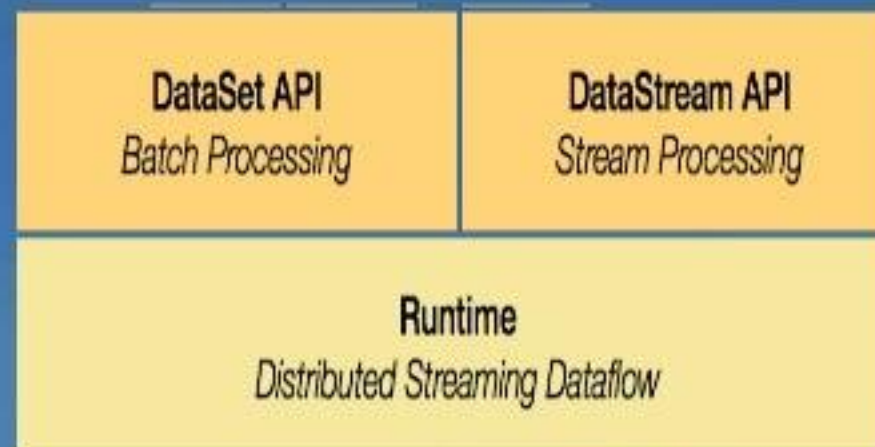
Item	Item			
	W	X	Y	Z
W	1.5	1.2	1.0	0.8
X	1.7	0.6	1.1	0.4
Y				
Z				

Item Matrix

Machine Learning at scale







Batch Processing



Graph Analysis

Why Apache Flink is Next-Gen?

			
<ul style="list-style-type: none"> Batch 	<ul style="list-style-type: none"> Batch Interactive 	<ul style="list-style-type: none"> Batch Interactive Near-Real time Streaming Iterative processing 	<ul style="list-style-type: none"> Batch Interactive Real-Time Streaming Native Iterative processing
MapReduce	Direct Acyclic Graphs (DAG) Dataflows	RDD: Resilient Distributed Datasets	Cyclic Dataflows
1 st Generation (1G)	2 nd Generation (2G)	3 rd Generation (3G)	4 th Generation (4G)

5. Flink APIs

5.1 **DataSet API** for **static data** - Java, Scala, and Python

5.2 **DataStream API** for **unbounded real-time streams** - Java and Scala

5.3 **Table API** for **relational queries** - Java and Scala

5.1 DataSet API – Batch processing

```
case class Word (word: String, frequency: Int)
```

DataSet API (batch):

```
val lines: DataSet[String] = env.readTextFile(...)
lines.flatMap {line => line.split(" ")}
      .map(word => Word(word, 1))
      .groupBy("word") .sum("frequency")
      .print()
```

DataStream API (streaming):

```
val lines: DataStream[String] = env.fromSocketStream(...)

lines.flatMap {line => line.split(" ")}
      .map(word => Word(word, 1))
      .window(Time.of(5, SECONDS)) .every(Time.of(1, SECONDS))
      .groupBy("word") .sum("frequency")
      .print()
```


5.2 DataStream API – Real-Time Streaming Analytics

➤ Many time-critical applications need to process **large streams of live data** and provide results in **real-time**. For example:

- Fraud detection
- Financial Stock monitoring
- Anomaly detection
- Traffic management applications
- Online recommenders

➤ **Flink Streaming** provides high-throughput, low-latency stateful stream processing system with rich windowing semantics. It has built-in **connectors** to many data sources like Flume, Kafka, Twitter, RabbitMQ

5.2 DataStream API – Real-Time Streaming Analytics

- Still in **Beta** as of June 24th 2015 (Flink 0.9 release)
- **Data streams** can be transformed and modified using high-level functions similar to the ones provided by the batch processing API.
- Flink Streaming provides **native support for iterative stream processing**.
- Streaming **Fault-Tolerance** added in Flink 0.9 (released on June 24th , 2015) allows **Exactly-once processing delivery guarantees** for Flink streaming programs that analyze streaming sources persisted by Apache Kafka. See paper: 'Lightweight Asynchronous Snapshots for Distributed Dataflows' <http://arxiv.org/pdf/1506.08603v1.pdf>

5.2 DataStream API – Real-Time Streaming Analytics

➤ Data Streaming Fault Tolerance document:

http://ci.apache.org/projects/flink/flink-docs-master/internals/stream_checkpointing.html

➤ Flink being based on a **pipelined execution engine** akin to parallel database systems allows:

- to integrate **streaming** operations with **rich windowing semantics** seamlessly
- process streaming operations in a pipelined way with **lower latency than micro-batch architectures** and **without the complexity of lambda architectures**.

➤ Flink Streaming web resources at the **Flink Knowledge**

Base <http://sparkbigdata.com/component/tags/tag/49-flink-streaming>

5.3 Table API – Relational Queries

Table API (queries)

```
val customers = envreadCsvFile(...).as('id, 'mktSegment)
    .filter("mktSegment = AUTOMOBILE")

val orders = env.readCsvFile(...)
    .filter( o =>
dateFormat.parse(o.orderDate).before(date) )
    .as("orderId, custId, orderDate, shipPrio")

val items = orders
    .join(customers).where("custId = id")
    .join(lineitems).where("orderId = id")
    .select("orderId, orderDate, shipPrio,
        extdPrice * (Literal(1.0f) - discount) as
revenue")

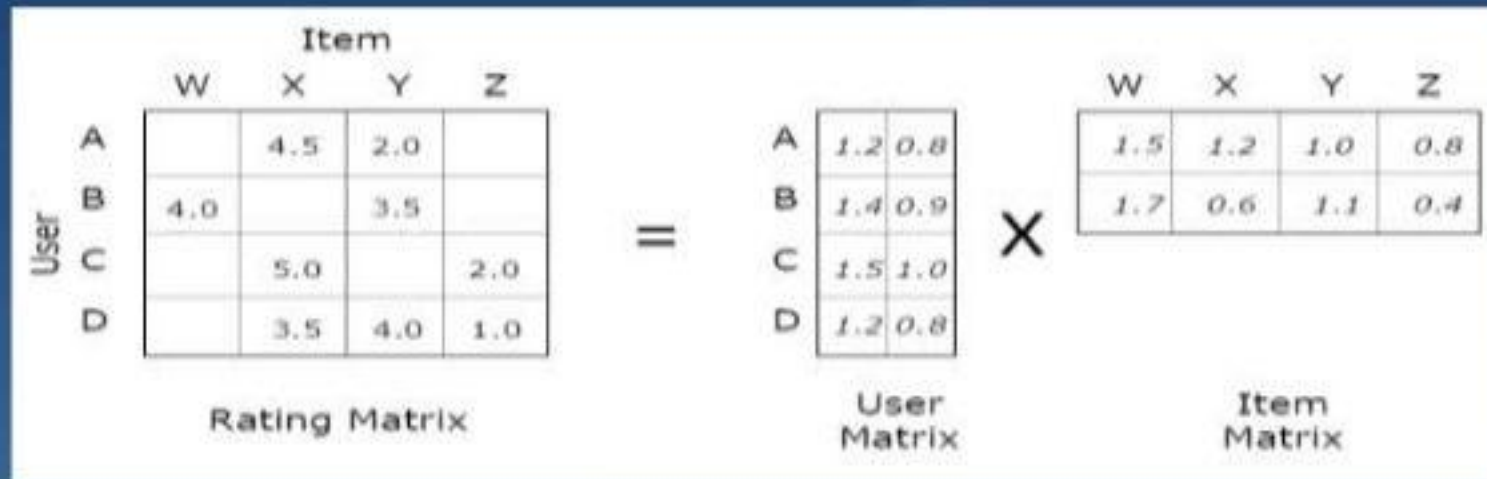
val result = items
    .groupBy("orderId, orderDate, shipPrio")
    .select("orderId, revenue.sum, orderDate, shipPrio")
```


5.3 Table API – Relational Queries

- **Table API** added in February 2015. Still in **Beta** as of June 24th 2015 (Flink 0.9 release)
- Flink provides Table API that allows specifying operations using **SQL-like expressions** instead of manipulating DataSet or DataStream.
- Table API can be used in both **batch** (on structured data sets) and **streaming** programs (on structured data streams). <http://ci.apache.org/projects/flink/flink-docs-master/libs/table.html>
- Flink Table web resources at the **Apache Flink Knowledge Base**: <http://sparkbigdata.com/component/tags/tag/52-flink-table>

6. Flink Domain Specific Libraries

6.1 FlinkML – Machine Learning Library



6.2 Gelly – Graph Analytics for Flink

